OPE-001 - Define, monitor and analyze customer workload health KPIs

Solution components are monitored via CloudWatch dashboards so they can centralize health status information.

This section presents how to monitor and check the health of resources configured in the customer's AWS portal. It shows the collection of information and logs used to perform load and metric analysis, indicating the use of tools such as CloudWatch and the process of capturing logs of operational events for later analysis. This allows for the generation of alerts based on specific metrics that can inform us of the status of resources at a given time and that could lead to actions to resolve the situation.
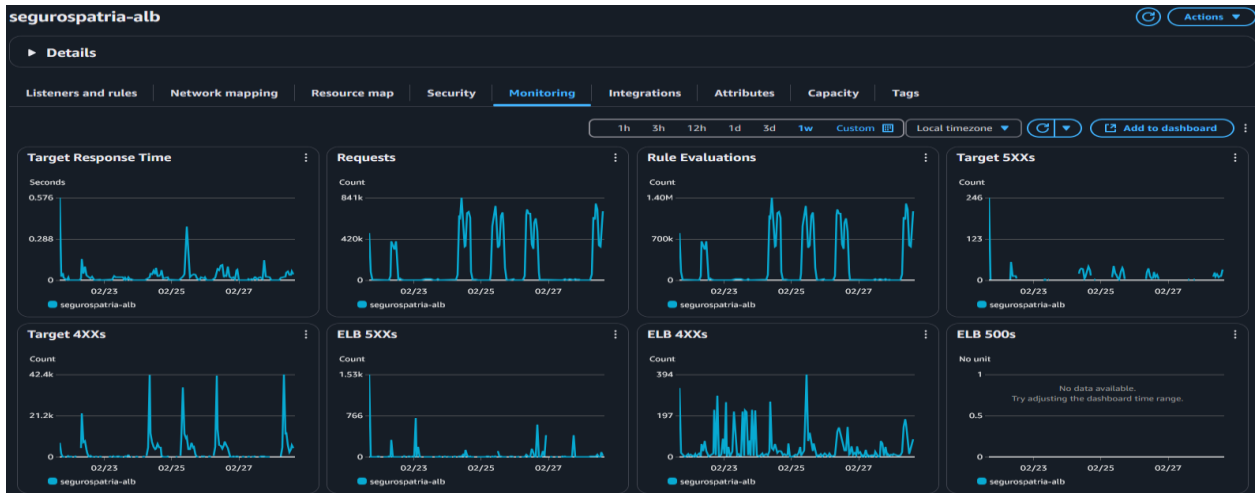
CloudWatch images for RDS performance:

Charts or DB metrics

As part of our monitoring practice, we always make sure to monitor important metrics such as CPU, RAM, and disk. In addition, we've configured advanced metrics and event viewer log forwarding to Cloudwatch through the CloudWatch agent.
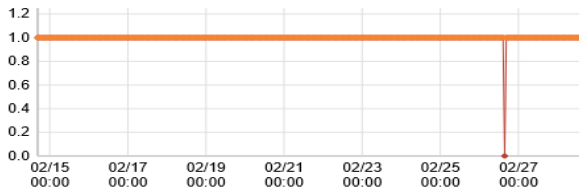
RDS Metrics

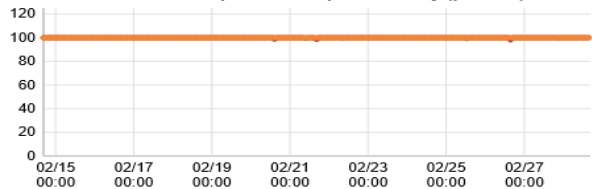## ALB Metrics

Metrics Route 53 Health Checks

Click on a graph to see an expanded view. View metrics in CloudWatch
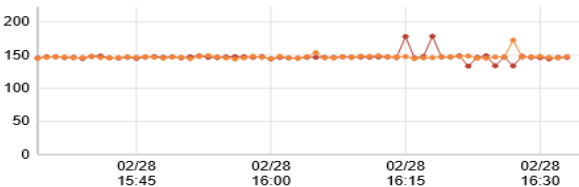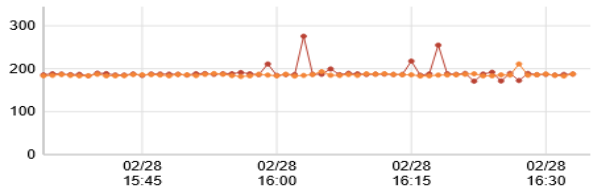
Auto Scaling Groups
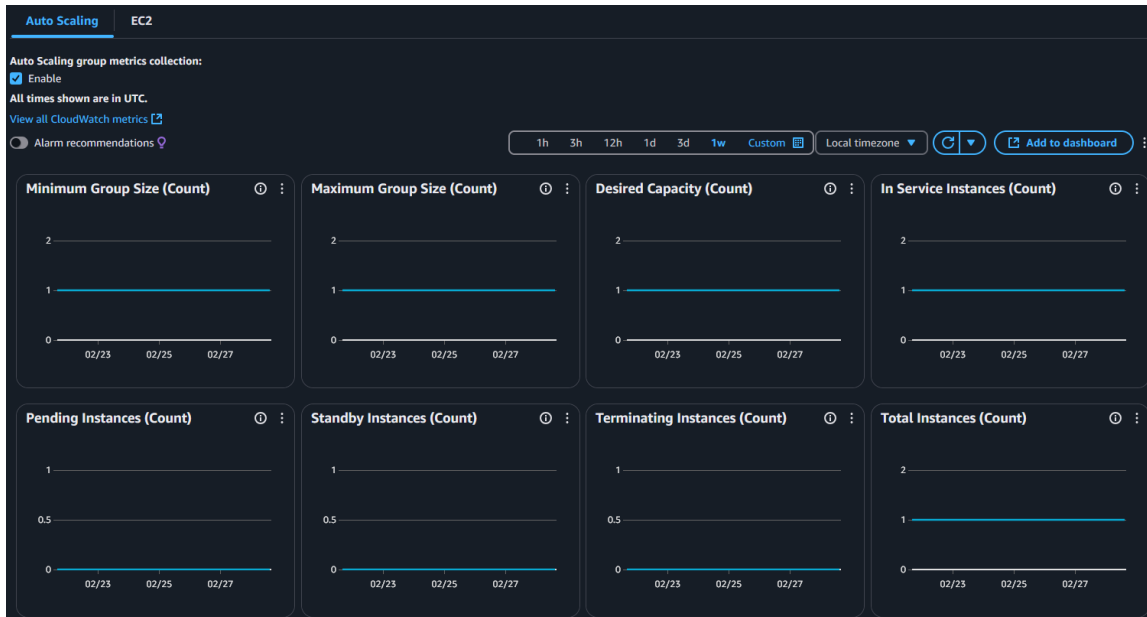
Monitoring alarms are defined:

Configured alarms are triggered in certain scenarios depending on the observed metric, for example: if the CPU percentage excedes 90%, it sends an alarm.

The main metrics monitored are CPU, RAM and disk.

Logging by CloudWatch

We have a number of Log Groups in which we store logs of different types, from RDS error logs to logs forwarded from the servers via the Cloudwatch agent installed on the servers.